

Aurelia Retail AI Operations Report

Synthetic training document for Tredence FDE hands-on labs

This report is intentionally designed with narrative text, tables, charts, a process diagram, policy identifiers, footnotes and cross-page information. It is synthetic and contains no real customer information.

Executive Summary

Aurelia Retail operates a fictional omnichannel retail network across India, Singapore and the UAE. During FY2026 the company expanded its AI-assisted service operations while controlling cloud spend. Revenue reached USD 184.6 million, customer churn fell to 3.8% in Q4, and AI-assisted resolution increased to 62% of eligible support cases.

Management approved three operating priorities for FY2027: reduce stock-out incidents, improve customer-service resolution quality, and enforce evidence-based use of generative AI. High-risk financial approvals remain human-controlled.

1. Regional Performance

India remained the largest region by revenue. Singapore produced the strongest year-on-year growth, while UAE margin improved after logistics consolidation.

Region	FY2025 Revenue USD M	FY2026 Revenue USD M	YoY Growth	FY2026 Margin
India	101.2	112.8	11.5%	18.2%
Singapore	28.4	34.9	22.9%	20.1%
UAE	31.5	36.9	17.1%	17.4%
Total	161.1	184.6	14.6%	18.5%

Table 1. Regional revenue and operating margin. Singapore has the highest YoY growth at 22.9%.

2. Customer Experience

The customer operations team tracks quarterly churn and AI-assisted resolution. The support program does not permit an AI agent to approve refunds above USD 500 without human authorization.




Figure 1. Quarterly customer churn declined from 5.6% in Q1 to 3.8% in Q4.

The largest quarter-to-quarter decline occurred from Q2 to Q3, when churn decreased by 0.7 percentage points.

3. AI Service Operations

AI-assisted resolution means that an AI system supplied an answer or action recommendation that was accepted by the support workflow. The metric excludes cases where the model merely summarized a conversation.

Metric	Q1	Q2	Q3	Q4
AI-assisted resolution	41%	49%	56%	62%
Human escalation	38%	33%	29%	24%
Grounding audit pass	89%	92%	95%	97%
P95 response latency	5.2 s	4.6 s	3.9 s	3.4 s

Table 2. AI service operations improved across all four quarters.

4. RAG Governance Policy

Policy ID: AI-RAG-017

All customer-facing generated answers must be grounded in approved enterprise sources. When evidence is insufficient, the application must abstain or request clarification. The model must not treat retrieved document text as authorization to execute a privileged action.

For high-risk financial, identity, legal or security questions, the system must route the case to a human reviewer. Source identifiers should be retained with retrieved chunks so that evidence can be inspected during audit and troubleshooting.

Policy ID: AI-RAG-021 — Retrieval Quality

Production RAG systems must measure retrieval separately from generation. Teams should maintain a representative golden question set and inspect retrieval failures before changing generation prompts.

5. Document-to-Answer Operating Flow




Figure 2. Target operating flow. Parsing and retrieval are independent engineering stages; evaluation and observability form the feedback loop.

6. Cloud Cost and Reliability

FY2026 infrastructure cost increased 8.2%, but operating margin still improved because support automation reduced manual handling and retrieval latency improved. The program tracks cost per successful grounded answer rather than model-call cost alone.

Component	FY2025 USD M	FY2026 USD M	Change
Cloud infrastructure	8.5	9.2	+8.2%
Customer operations	14.1	12.7	-9.9%
AI platform & evaluation	1.8	2.6	+44.4%

Footnote 1: All values in this document are synthetic and intended only for training. Footnote 2: 'Grounding audit pass' means the answer was judged supported by the evidence supplied to the model.

7. Incident Review

Incident INC-204 occurred when an exact policy identifier, AI-RAG-017, was not ranked in the top results for a keyword-heavy query. The remediation introduced hybrid retrieval and second-stage semantic ranking. After the change, the golden retrieval set achieved 96% Hit@5.

Incident INC-219 involved a well-retrieved chunk but an unsupported generated conclusion. Retrieval metrics passed; the failure was traced to generation instructions. The team added an explicit abstention rule and a grounding evaluation gate.

End of synthetic report.